

```
In [ ]: import marimo as mo
```

VQ-VAE Tokenizer Evaluation

Protein / Ligand VQ-VAE の学習後の定量的性能評価ノートブック。

評価項目:

1. 再構成誤差 (MSE) — 全体 & 次元別
2. Codebook 利用率 & Perplexity
3. Original vs Reconstructed の散布図
4. Codebook 使用頻度の分布
5. 潜在空間の可視化 (t-SNE)

```
In [ ]: import sys
from pathlib import Path

import matplotlib.pyplot as plt
import numpy as np
import torch

# Add project root to path (resolve from this file, not cwd)
project_root = Path(__file__).resolve().parent.parent
if str(project_root) not in sys.path:
    sys.path.insert(0, str(project_root))

from src.config import CrossDockedConfig, VQVAETrainingConfig
from src.data.descriptors import ComplexDescriptorDataModule
from src.model.vqvae_module import VQVAEModule

# '%matplotlib inline' command supported automatically in marimo
plt.rcParams["figure.dpi"] = 120
```

1. Load checkpoint and data

```
In [ ]: # --- Checkpoint path (edit here) ---
# Best checkpoint is typically the one with lowest val/protein_recon.
# List available checkpoints:
ckpt_dir = project_root / "checkpoints"
if ckpt_dir.exists():
    ckpts = sorted(ckpt_dir.rglob("*.ckpt"))
    for _p in ckpts:
        print(_p.relative_to(project_root))
else:
    print("No checkpoints/ directory found. Check wandb run directory.")
    ckpts = sorted(project_root.rglob("*.ckpt"))
    for _p in ckpts[-5:]:
        print(_p.relative_to(project_root))
```

No checkpoints/ directory found. Check wandb run directory.
 pocket-ligand-vqvae/uhhyc6y7/checkpoints/vqvae-epoch=86-val/protein_recon=0.1287.ckpt
 pocket-ligand-vqvae/uhhyc6y7/checkpoints/vqvae-epoch=93-val/protein_recon=0.1288.ckpt
 pocket-ligand-vqvae/yggua4f0/checkpoints/vqvae-epoch=90-val/protein_recon=0.1316.ckpt
 pocket-ligand-vqvae/yggua4f0/checkpoints/vqvae-epoch=95-val/protein_recon=0.1316.ckpt
 pocket-ligand-vqvae/yggua4f0/checkpoints/vqvae-epoch=99-val/protein_recon=0.1316.ckpt

```
In [ ]: import os

# Override via `VQVAE_CKPT` env var so a single notebook can be exported
# per run without editing source in between. Falls back to the baseline
# (yggua4f0) used in the original evaluation.
CKPT_PATH = os.environ.get(
    "VQVAE_CKPT",
    "/gs/bs/tga-ohuelab/sakano/git/pocket-conditioned-ligand-gen/pocket-ligand-vqvae/ql4ls8h6/checkpoints/vqvae-epoch=96-val/protein_recon=0.1238.ckpt"
)

print(f>Loading: {CKPT_PATH}")

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
module = VQVAEModule.load_from_checkpoint(str(CKPT_PATH), map_location=device)
module.eval()
module.to(device)

protein_vqvae = module.protein_vqvae
ligand_vqvae = module.ligand_vqvae
print(f>Device: {device}")
```

Loading: /gs/bs/tga-ohuelab/sakano/git/pocket-conditioned-ligand-gen/pocket-ligand-vqvae/ql4ls8h6/checkpoints/vqvae-epoch=96-val/protein_recon=0.1238.ckpt

/gs/bs/tga-ohuelab/sakano/git/pocket-conditioned-ligand-gen/.venv/lib/python3.12/site-packages/torch/nn/modules/transformer.py:392: UserWarning: enable_nested_tensor is True, but self.use_nested_tensor is False because encoder_layer.norm_first was True

warnings.warn(

Device: cuda

```
In [ ]: # Load cached descriptors and prepare test split
import os as _os
from pathlib import Path as _Path

config = VQVAETrainingConfig()
data_config = CrossDockedConfig(data_dir=project_root / "data")
dm = ComplexDescriptorDataModule(config, data_config)
# Override the cache dir so B2 (v2 cache + skip_sincos stats) can be
# evaluated against the matching shards/stats. Defaults to the v1 cache.
_cache_override = _os.environ.get("VQVAE_CACHE_DIR")
if _cache_override:
    dm.cache_dir = _Path(_cache_override)
dm.setup()
```

```

norm_stats = dm.norm_stats

# The full test split is ~2M complexes across 436 shards; loading it all
# takes the better part of an hour and is overkill for MSE / t-SNE
# statistics. Cap the notebook at MAX_TEST_COMPLEXES by stopping early.
MAX_TEST_COMPLEXES = 2000

if dm.protein_test is None:
    prot_mean_np = norm_stats["protein_mean"].numpy()
    prot_std_np = norm_stats["protein_std"].numpy()
    lig_mean_np = norm_stats["ligand_mean"].numpy()
    lig_std_np = norm_stats["ligand_std"].numpy()

    protein_test_pockets = []
    ligand_test_molecules = []
    # Read each test shard once and pull out both protein + ligand,
    # so we don't pay 2x disk I/O across two separate iterators.
    for shard_idx, local_indices in dm._test_plan:
        if len(protein_test_pockets) >= MAX_TEST_COMPLEXES:
            break
        shard_path = dm._shard_dir / f"shard_{shard_idx:04d}.pt"
        shard_data = torch.load(shard_path, weights_only=False)
        for local_idx in local_indices:
            if len(protein_test_pockets) >= MAX_TEST_COMPLEXES:
                break
            cplx = shard_data[local_idx]
            protein_test_pockets.append(
                torch.from_numpy(
                    (cplx["protein"] - prot_mean_np) / prot_std_np,
                )
                .float()
                .to(device),
            )
            ligand_test_molecules.append(
                torch.from_numpy(
                    (cplx["ligand"] - lig_mean_np) / lig_std_np,
                )
                .float()
                .to(device),
            )
        del shard_data
    else:
        protein_test_pockets = [
            t.to(device) for t in dm.protein_test[:MAX_TEST_COMPLEXES]
        ]
        ligand_test_molecules = [
            t.to(device) for t in dm.ligand_test[:MAX_TEST_COMPLEXES]
        ]

    protein_test_flat = torch.cat(protein_test_pockets)
    ligand_test_flat = torch.cat(ligand_test_molecules)

    print(
        f"Protein test: {len(protein_test_pockets)} pockets, "
        f"{protein_test_flat.shape[0]} residues total"
    )

```

```
)
print(
    f"Ligand test: {len(ligand_test_molecules)} molecules, "
    f"{ligand_test_flat.shape[0]} atoms total"
)
```

Protein test: 2000 pockets, 53813 residues total

Ligand test: 2000 molecules, 55320 atoms total

```
In [ ]: # Both protein and ligand share the same TransformerVQVAE architecture and
# consume variable-length per-pocket / per-molecule sequences.
@torch.no_grad()
def run_vqvae(model: "TransformerVQVAE", sequences: "list[Tensor]"):
    """Run TransformerVQVAE per-sequence and collect recon/indices/z."""
    all_recon, all_indices, all_z = [], [], []
    for seq in sequences:
        if seq.shape[0] == 0:
            continue
        x_seq = seq.unsqueeze(0) # (1, N, D)
        h = model.input_proj(model.input_norm(x_seq))
        h = h + model.pos_encoding[: seq.shape[0]]
        h = model.transformer_encoder(h)
        z = model.latent_norm(model.latent_proj(h)).squeeze(0) # (N, latent_dim)
        quantized, indices, _ = model.codebook(z)
        q_seq = model.latent_unproj(quantized).unsqueeze(0)
        dec_in = q_seq + model.pos_encoding[: seq.shape[0]]
        dec_out = model.transformer_decoder(dec_in)
        recon = model.output_proj(dec_out).squeeze(0) # (N, D)
        all_recon.append(recon)
        all_indices.append(indices)
        all_z.append(z)
    return torch.cat(all_recon), torch.cat(all_indices), torch.cat(all_z)

prot_recon, prot_indices, prot_z = run_vqvae(protein_vqvae, protein_test_pockets)
lig_recon, lig_indices, lig_z = run_vqvae(ligand_vqvae, ligand_test_molecules)

print("Inference done.")
print(
    f"Protein: recon {prot_recon.shape}, indices {prot_indices.shape}, z {prot_z.shape}"
)
print(
    f"Ligand: recon {lig_recon.shape}, indices {lig_indices.shape}, z {lig_z.shape}"
)
```

Inference done.

Protein: recon torch.Size([53813, 12]), indices torch.Size([53813]), z torch.Size([53813, 16])

Ligand: recon torch.Size([55320, 4]), indices torch.Size([55320]), z torch.Size([55320, 8])

2. Reconstruction Error (MSE)

正規化空間と元スケールの両方で再構成誤差を評価する。

```

In [ ]: def compute_mse_metrics(
        original: "Tensor",
        reconstructed: "Tensor",
        name: str,
        norm_mean: "Tensor",
        norm_std: "Tensor",
    ):
        """Compute MSE in both normalized and original scale."""
        orig_np = original.cpu().numpy()
        recon_np = reconstructed.cpu().numpy()

        # Normalized space
        mse_norm = np.mean((orig_np - recon_np) ** 2)
        per_dim_mse_norm = np.mean((orig_np - recon_np) ** 2, axis=0)

        # De-normalize to original scale
        mean_np = norm_mean.numpy()
        std_np = norm_std.numpy()
        orig_denorm = orig_np * std_np + mean_np
        recon_denorm = recon_np * std_np + mean_np
        mse_orig = np.mean((orig_denorm - recon_denorm) ** 2)
        per_dim_mse_orig = np.mean((orig_denorm - recon_denorm) ** 2, axis=0)

        print(f"=== {name} VQ-VAE ===")
        print(f" Overall MSE (normalized): {mse_norm:.6f}")
        print(f" Overall MSE (original): {mse_orig:.6f}")
        print(f" Overall RMSE (original): {np.sqrt(mse_orig):.6f}")
        print()

        return {
            "mse_norm": mse_norm,
            "per_dim_mse_norm": per_dim_mse_norm,
            "mse_orig": mse_orig,
            "per_dim_mse_orig": per_dim_mse_orig,
            "orig_denorm": orig_denorm,
            "recon_denorm": recon_denorm,
        }

    prot_metrics = compute_mse_metrics(
        protein_test_flat,
        prot_recon,
        "Protein",
        norm_stats["protein_mean"],
        norm_stats["protein_std"],
    )
    lig_metrics = compute_mse_metrics(
        ligand_test_flat,
        lig_recon,
        "Ligand",
        norm_stats["ligand_mean"],
        norm_stats["ligand_std"],
    )

```

```

=== Protein VQ-VAE ===
Overall MSE (normalized):  0.130986
Overall MSE (original):    0.084838
Overall RMSE (original):   0.291270

=== Ligand VQ-VAE ===
Overall MSE (normalized):  0.017311
Overall MSE (original):    0.010141
Overall RMSE (original):   0.100704

```

2.1 Per-dimension MSE

Protein: 12-D backbone Z-matrix (4 values $\times$ 3 atoms N/CA/C). 4 values per atom are (bond_length, bond_angle, sin_torsion, cos_torsion) for continuation residues, or pocket-frame-anchored spherical coords for segment-start residues. Ligand: dim 0 = bond length, dim 1 = bond angle, dim 2-3 = sin/cos dihedral

```


In [ ]: _fig, _axes = plt.subplots(2, 2, figsize=(14, 8))
        prot_dim_labels = [
            "N_d",
            "N_theta",
            "N_sin",
            "N_cos",
            "CA_d",
            "CA_theta",
            "CA_sin",
            "CA_cos",
            "C_d",
            "C_theta",
            "C_sin",
            "C_cos",
        ]
        n_prot_dims = len(prot_dim_labels)
        # Protein - normalized
        _axes[0, 0].bar(range(n_prot_dims), prot_metrics["per_dim_mse_norm"])
        _axes[0, 0].set_xticks(range(n_prot_dims), prot_dim_labels, rotation=90, for
        _axes[0, 0].set_title("Protein - Per-dim MSE (normalized)")
        _axes[0, 0].set_ylabel("MSE")
        _axes[0, 1].bar(range(n_prot_dims), prot_metrics["per_dim_mse_orig"])
        _axes[0, 1].set_xticks(range(n_prot_dims), prot_dim_labels, rotation=90, for
        # Protein - original scale
        _axes[0, 1].set_title("Protein - Per-dim MSE (original scale)")
        _axes[0, 1].set_ylabel("MSE")
        lig_dim_labels = ["bond_len", "bond_angle", "sin_dih", "cos_dih"]
        _axes[1, 0].bar(range(4), lig_metrics["per_dim_mse_norm"])
        _axes[1, 0].set_xticks(range(4), lig_dim_labels, rotation=90, fontsize=7)
        # Ligand - normalized
        _axes[1, 0].set_title("Ligand - Per-dim MSE (normalized)")
        _axes[1, 0].set_ylabel("MSE")
        _axes[1, 1].bar(range(4), lig_metrics["per_dim_mse_orig"])
        _axes[1, 1].set_xticks(range(4), lig_dim_labels, rotation=90, fontsize=7)
        _axes[1, 1].set_title("Ligand - Per-dim MSE (original scale)")

```

```

_axes[1, 1].set_ylabel("MSE")
# Ligand - original scale
_fig.tight_layout()
_fig

```



3. Codebook Utilization & Perplexity

理想的なcodebookは全コードが均等に使われる。

- **Utilization:** 使用されたコード数 / 全コード数 (1.0が理想)
- **Perplexity:** $\exp(\text{entropy})$ 。均等利用時は codebook_size と一致

```

In [ ]: def codebook_stats(indices: "Tensor", codebook_size: int, name: str):
        """Compute and print codebook utilization metrics."""
        idx_np = indices.cpu().numpy()
        unique = np.unique(idx_np)
        utilization = len(unique) / codebook_size

        counts = np.bincount(idx_np, minlength=codebook_size).astype(float)
        probs = counts / counts.sum()
        probs_nonzero = probs[probs > 0]
        entropy = -np.sum(probs_nonzero * np.log(probs_nonzero))
        perplexity = np.exp(entropy)

        # Max perplexity = codebook_size (uniform distribution)
        max_perplexity = codebook_size
        normalized_perplexity = perplexity / max_perplexity

        print(f"=== {name} Codebook ===")
        print(f"  Codebook size:      {codebook_size}")
        print(f"  Active codes:        {len(unique)}")
        print(f"  Utilization:         {utilization:.4f}")
        print(f"  Perplexity:          {perplexity:.1f} / {max_perplexity}")

```

```

    print(f"   Normalized perplexity: {normalized_perplexity:.4f}")
    print(f"   Dead codes:                {codebook_size - len(unique)}")
    print()
    return counts

prot_counts = codebook_stats(prot_indices, config.protein.codebook_size, "Protein")
lig_counts = codebook_stats(lig_indices, config.ligand.codebook_size, "Ligand")

=== Protein Codebook ===
Codebook size:      2048
Active codes:       2028
Utilization:        0.9902
Perplexity:         1594.4 / 2048
Normalized perplexity: 0.7785
Dead codes:         20

=== Ligand Codebook ===
Codebook size:      1024
Active codes:       1024
Utilization:        1.0000
Perplexity:         763.1 / 1024
Normalized perplexity: 0.7452
Dead codes:         0


```

3.1 Codebook usage distribution

```

In [ ]: _fig, _axes = plt.subplots(1, 2, figsize=(14, 4))
prot_sorted = np.sort(prot_counts[::-1])
# Protein codebook usage – sorted descending
_axes[0].bar(range(len(prot_sorted)), prot_sorted, width=1.0)
_axes[0].set_xlabel("Code rank")
_axes[0].set_ylabel("Count")
_axes[0].set_title(f"Protein Codebook Usage (size={config.protein.codebook_size})")
_axes[0].set_yscale("log")
lig_sorted = np.sort(lig_counts[::-1])
_axes[1].bar(range(len(lig_sorted)), lig_sorted, width=1.0)
# Ligand codebook usage – sorted descending
_axes[1].set_xlabel("Code rank")
_axes[1].set_ylabel("Count")
_axes[1].set_title(f"Ligand Codebook Usage (size={config.ligand.codebook_size})")
_axes[1].set_yscale("log")
_fig.tight_layout()
_fig

```



4. Original vs Reconstructed scatter plots

各次元で元の値 vs 再構成値をプロットし、対角線からのずれを確認する。元スケールで表示。

```
In [ ]: def scatter_orig_vs_recon(
    orig: "ndarray",
    recon: "ndarray",
    dim_labels: list[str],
    title: str,
    max_points: int = 5000,
):
    """Scatter plot of original vs reconstructed for selected dimensions."""
    n_dims = orig.shape[1]
    max_inline_dims = 8
    if n_dims <= max_inline_dims:
        selected = list(range(n_dims))
    else:
        selected = [
            0,
            n_dims // 4,
            n_dims // 2,
            3 * n_dims // 4,
            n_dims - 2,
            n_dims - 1,
        ]
    n_plots = len(selected)
    _fig, _axes = plt.subplots(1, n_plots, figsize=(3.5 * n_plots, 3.2))
    if n_plots == 1:
        _axes = [_axes]
    rng = np.random.default_rng(42)
    idx = rng.choice(len(orig), min(max_points, len(orig)), replace=False)
    for ax, dim in zip(_axes, selected, strict=False):
        ax.scatter(orig[idx, dim], recon[idx, dim], s=1, alpha=0.3)
        lo = min(orig[idx, dim].min(), recon[idx, dim].min())
        hi = max(orig[idx, dim].max(), recon[idx, dim].max())
        ax.plot([lo, hi], [lo, hi], "r--", linewidth=0.8)
        ax.set_xlabel("Original")
        ax.set_ylabel("Reconstructed")
        ax.set_title(dim_labels[dim], fontsize=9)
        ax.set_aspect("equal", adjustable="datalim") # Subsample for plotti
    _fig.suptitle(title, fontsize=12)
    _fig.tight_layout()
    plt.show()


scatter_orig_vs_recon(
    prot_metrics["orig_denorm"],
    prot_metrics["recon_denorm"],
    prot_dim_labels,
    "Protein: Original vs Reconstructed (original scale)",
)

scatter_orig_vs_recon(
    lig_metrics["orig_denorm"],
```

```

lig_metrics["recon_denorm"],
lig_dim_labels,
"Ligand: Original vs Reconstructed (original scale)",
)

```



5. Reconstruction error distribution

サンプルごとの再構成誤差の分布を確認し、外れ値の存在を把握する。

```

In [ ]: _fig, _axes = plt.subplots(1, 2, figsize=(12, 4))
prot_per_sample = np.mean(
    (protein_test_flat.cpu().numpy() - prot_recon.cpu().numpy()) ** 2, axis=
)
lig_per_sample = np.mean(
    (ligand_test_flat.cpu().numpy() - lig_recon.cpu().numpy()) ** 2, axis=1
)
_axes[0].hist(prot_per_sample, bins=100, edgecolor="none", alpha=0.8)
_axes[0].axvline(
    np.median(prot_per_sample),
    color="r",
    linestyle="--",
    label=f"median={np.median(prot_per_sample):.4f}",
)
_axes[0].axvline(
    np.mean(prot_per_sample),
    color="orange",
    linestyle="--",
    label=f"mean={np.mean(prot_per_sample):.4f}",
)
_axes[0].set_xlabel("Per-sample MSE")
_axes[0].set_ylabel("Count")
_axes[0].set_title("Protein - Per-sample MSE distribution")
_axes[0].legend(fontsize=8)
_axes[1].hist(lig_per_sample, bins=100, edgecolor="none", alpha=0.8)
_axes[1].axvline(
    np.median(lig_per_sample),
    color="r",
    linestyle="--",
    label=f"median={np.median(lig_per_sample):.4f}",
)

```

```

_axes[1].axvline(
    np.mean(lig_per_sample),
    color="orange",
    linestyle="--",
    label=f"mean={np.mean(lig_per_sample):.4f}",
)
_axes[1].set_xlabel("Per-sample MSE")
_axes[1].set_ylabel("Count")
_axes[1].set_title("Ligand - Per-sample MSE distribution")
_axes[1].legend(fontsize=8)
_fig.tight_layout()
for name, mse_arr in [("Protein", prot_per_sample), ("Ligand", lig_per_sample)]:
    print(f"{name} per-sample MSE percentiles:")
    for _p in [50, 75, 90, 95, 99]:
        print(f"  P{_p}: {np.percentile(mse_arr, _p):.6f}")
    print()
_fig


```

Protein per-sample MSE percentiles:

P50: 0.085266
 P75: 0.179838
 P90: 0.294688
 P95: 0.385755
 P99: 0.640828

Ligand per-sample MSE percentiles:

P50: 0.001130
 P75: 0.003083
 P90: 0.008802
 P95: 0.022587
 P99: 0.342676



5.5 3D Reconstruction RMSD (Å)

記述子空間での再構成誤差ではなく、実際の3次元座標に復元した上での RMSD を評価する。

- **Protein:** backbone (N, CA, C) 座標の RMSD
- **Ligand:** 重原子座標の RMSD

2 種類の RMSD を並べて報告する:

- **Per-atom (superposition なし):** pocket frame を共有しているため、segment-start のズレや NeRF 累積による剛体シフトがそのまま出る。生成 pipeline 全体の精度。
- **Kabsch-aligned:** 原子集合を剛体 (回転 + 並進) でベストフィットさせた後の RMSD。内部変形 (conformer shape) の再現精度のみを反映する。

複数の複合体をサンプリングし、VQ-VAE の encode → decode → 記述子逆変換 → 3D座標 のパイプライン全体の精度を確認する。

```
In [ ]: import pyarrow.parquet as pq

from src.config import PocketExtractionConfig
from src.tokenizers.ligand import LigandDescriptor, parse_sdf
from src.tokenizers.protein import (
    BackboneZMatrixDescriptor,
    _compute_canonical_frame,
    extract_pocket,
)

def kabsch_rmsd(p: np.ndarray, q: np.ndarray) -> float:
    """Per-atom RMSD after optimal rigid-body alignment of q onto p (Kabsch)

    Both inputs shape (N, 3). Returns RMSD in the same units.
    """
    p_c = p - p.mean(axis=0)
    q_c = q - q.mean(axis=0)
    h = q_c.T @ p_c
    u, _, vt = np.linalg.svd(h)
    d = np.sign(np.linalg.det(vt.T @ u.T))
    rot = vt.T @ np.diag([1.0, 1.0, d]) @ u.T
    q_aligned = q_c @ rot.T
    return float(np.sqrt(np.mean(np.sum((p_c - q_aligned) ** 2, axis=-1))))

pocket_config = PocketExtractionConfig()
protein_desc_calc = BackboneZMatrixDescriptor()
ligand_desc_calc = LigandDescriptor()
# Use the HuggingFace hub cache manifest instead of the CrossDocked2020
# types/ tree (which isn't present on this machine).
hub_cache_dir = project_root / "data" / "hub_cache"
manifest_path = hub_cache_dir / "repo" / "manifest.parquet"
receptor_dir = hub_cache_dir / "receptors"
ligand_dir = hub_cache_dir / "ligands"
manifest_df = pq.read_table(manifest_path).to_pandas()
# Filter to the cdonly test split for fold 0 – matches the original
# cdonly*_test0.types file that this cell used to read.
test_df = manifest_df[
    (manifest_df["source_type"] == "cdonly")
    & (manifest_df["cdonly_fold0"] == "test")
].reset_index(drop=True)
entries = [
    (
        f"{row.complex_dir}/{row.receptor_pdb}",
        f"{row.pair_idx:07d}.sdf.gz",
    )
]
```

```

        0,
    )
    for row in test_df.itertuples(index=False)
]
prot_mean = norm_stats["protein_mean"].to(device)
prot_std = norm_stats["protein_std"].to(device)
lig_mean = norm_stats["ligand_mean"].to(device)
lig_std = norm_stats["ligand_std"].to(device)
N_SAMPLES = 200
rng = np.random.default_rng(42)
sample_indices = rng.choice(
    len(entries), size=min(N_SAMPLES * 5, len(entries)), replace=False
)
prot_rmsd_list = []
prot_rmsd_aligned_list = []
lig_rmsd_list = []
lig_rmsd_aligned_list = []
n_done = 0
for idx in sample_indices:
    if n_done >= N_SAMPLES:
        break
    rec_rel, lig_rel, _pose_idx = entries[idx]
    rec_path = receptor_dir / rec_rel
    lig_path = ligand_dir / lig_rel
    if not rec_path.exists() or not lig_path.exists():
        continue
    try:
        molecules = parse_sdf(lig_path)
        if not molecules:
            continue
        mol = molecules[0]
        lig_coords_orig = np.array(
            [(a[1], a[2], a[3]) for a in mol["atoms"] if a[0] != "H"],
            dtype=np.float32,
        )
        pocket_result = extract_pocket(rec_path, lig_coords_orig, pocket_cor
        if pocket_result is None:
            continue
        backbone_coords_orig, _seq, residue_ids = pocket_result
        ca_coords = backbone_coords_orig[:, 1].astype(np.float64)
        centroid, rotation = _compute_canonical_frame(ca_coords)
        pocket_frame = (centroid, rotation)
        prot_desc, prot_meta = protein_desc_calc.compute(
            backbone_coords_orig,
            residue_ids,
            pocket_frame=pocket_frame,
        )
        prot_t = torch.from_numpy(prot_desc).to(device)
        prot_norm = (prot_t - prot_mean) / prot_std
        with torch.no_grad():
            pi = protein_vqvae.encode(prot_norm)
            prot_recon_norm = protein_vqvae.decode(pi)
        prot_recon_desc = (prot_recon_norm * prot_std + prot_mean).cpu().num
        backbone_recon = BackboneZMatrixDescriptor.descriptor_to_backbone_cc
        prot_recon_desc, prot_meta
    )

```

```

prot_rmsd = np.sqrt(
    np.mean(np.sum((backbone_coords_orig - backbone_recon) ** 2, axis=1))
)
prot_rmsd_list.append(prot_rmsd)
prot_flat_orig = backbone_coords_orig.reshape(-1, 3).astype(np.float64)
prot_flat_recon = backbone_recon.reshape(-1, 3).astype(np.float64)
prot_rmsd_aligned_list.append(
    kabsch_rmsd(prot_flat_orig, prot_flat_recon)
)
lig_desc, _elements, lig_meta = ligand_desc_calc.compute(
    mol["atoms"], mol["bonds"], pocket_frame=pocket_frame
)
if len(lig_desc) == 0:
    continue
lig_t = torch.from_numpy(lig_desc).to(device)
lig_norm = (lig_t - lig_mean) / lig_std
with torch.no_grad():
    li = ligand_vqvae.encode(lig_norm)
    lig_recon_norm = ligand_vqvae.decode(li)
lig_recon_desc = (lig_recon_norm * lig_std + lig_mean).cpu().numpy()
lig_coords_recon = LigandDescriptor.descriptor_to_coords(
    lig_recon_desc, lig_meta, pocket_frame=pocket_frame
)
heavy_atoms = [a for a in mol["atoms"] if a[0] != "H"]
lig_coords_orig_arr = np.array(
    [(a[1], a[2], a[3]) for a in heavy_atoms], dtype=np.float64
)
lig_rmsd = np.sqrt(
    np.mean(
        np.sum((lig_coords_orig_arr - lig_coords_recon) ** 2, axis=-1)
    )
)
lig_rmsd_list.append(lig_rmsd)
lig_rmsd_aligned_list.append(
    kabsch_rmsd(lig_coords_orig_arr, lig_coords_recon.astype(np.float64))
)
n_done += 1
except Exception: # noqa: BLE001, S112
    continue
prot_rmsd_arr = np.array(prot_rmsd_list)
prot_rmsd_aligned_arr = np.array(prot_rmsd_aligned_list)
lig_rmsd_arr = np.array(lig_rmsd_list)
lig_rmsd_aligned_arr = np.array(lig_rmsd_aligned_list)
print(f"Evaluated {len(prot_rmsd_arr)} complexes")
print()

def _print_rmsd_stats(name: str, arr: np.ndarray) -> None:
    print(f"{name}:")
    print(f"  Mean:   {arr.mean():.4f}")
    print(f"  Median: {np.median(arr):.4f}")
    print(f"  Std:    {arr.std():.4f}")

_print_rmsd_stats("Protein backbone RMSD – per-atom (Å)", prot_rmsd_arr)
print()
_print_rmsd_stats(
    "Protein backbone RMSD – Kabsch-aligned (Å)", prot_rmsd_aligned_arr
)

```

```

)
print()
_print_rmsd_stats("Ligand heavy-atom RMSD – per-atom (Å)", lig_rmsd_arr)
print()
_print_rmsd_stats(
    "Ligand heavy-atom RMSD – Kabsch-aligned (Å)", lig_rmsd_aligned_arr
)

def _plot_rmsd_hist(ax, arr: np.ndarray, title: str) -> None:
    ax.hist(arr, bins=50, edgecolor="none", alpha=0.8)
    ax.axvline(
        np.median(arr),
        color="r",
        linestyle="--",
        label=f"median={np.median(arr):.3f} Å",
    )
    ax.axvline(
        arr.mean(),
        color="orange",
        linestyle="--",
        label=f"mean={arr.mean():.3f} Å",
    )
    ax.set_xlabel("RMSD (Å)")
    ax.set_ylabel("Count")
    ax.set_title(title)
    ax.legend(fontsize=8)

_fig, _axes = plt.subplots(2, 2, figsize=(12, 8))
_plot_rmsd_hist(_axes[0, 0], prot_rmsd_arr, "Protein backbone – per-atom")
_plot_rmsd_hist(
    _axes[0, 1], prot_rmsd_aligned_arr, "Protein backbone – Kabsch-aligned"
)
_plot_rmsd_hist(_axes[1, 0], lig_rmsd_arr, "Ligand heavy-atom – per-atom")
_plot_rmsd_hist(
    _axes[1, 1], lig_rmsd_aligned_arr, "Ligand heavy-atom – Kabsch-aligned"
)
_fig.tight_layout()
_fig

```

Evaluated 998 complexes

Protein backbone RMSD – per-atom (Å):

Mean: 1.4997

Median: 1.4692

Std: 0.5114

Protein backbone RMSD – Kabsch-aligned (Å):

Mean: 1.3740

Median: 1.3529

Std: 0.4673

Ligand heavy-atom RMSD – per-atom (Å):

Mean: 2.1796

Median: 1.4104


Std: 2.3056

Ligand heavy-atom RMSD – Kabsch-aligned (Å):

Mean: 0.2494

Median: 0.1619

Std: 0.2381



6. Latent space visualization (t-SNE)

Codebook ベクトルとエンコーダ出力を t-SNE で2次元に射影し、量子化の質を確認する。

```
In [ ]: from sklearn.manifold import TSNE

def plot_latent_tsne(
```



```


encoder_z: "Tensor",
codebook: "EMACodebook",
_codebook_size: int,
name: str,
n_samples: int = 5000,
) -> None:
    """t-SNE visualization of encoder outputs and codebook vectors.

    Vectors are plotted as-is; the codebook does plain L2 nearest-neighbor
    lookup on raw embeddings (encoder z is post-LayerNorm).
    """
    z_arr = encoder_z.cpu().numpy()
    cb_arr = codebook.embedding.cpu().detach().numpy()
    rng = np.random.default_rng(42)
    idx = rng.choice(len(z_arr), min(n_samples, len(z_arr)), replace=False)
    z_sub = z_arr[idx]
    combined = np.vstack([z_sub, cb_arr]) # Subsample encoder outputs
    tsne = TSNE(
        n_components=2, random_state=42, perplexity=min(30, len(combined)) //
    )
    embedded = tsne.fit_transform(combined)
    z_emb = embedded[: len(z_sub)]
    cb_emb = embedded[len(z_sub) :]
    _fig, ax = plt.subplots(figsize=(8, 6)) # Combine for t-SNE
    ax.scatter(
        z_emb[:, 0],
        z_emb[:, 1],
        s=1,
        alpha=0.2,
        c="steelblue",
        label="Encoder output",
    )
    ax.scatter(
        cb_emb[:, 0],
        cb_emb[:, 1],
        s=30,
        c="red",
        marker="x",
        linewidths=1,
        label="Codebook",
    )
    ax.set_title(f"{name} - t-SNE of latent space")
    ax.legend(fontsize=9)
    ax.set_xticks([])
    ax.set_yticks([])
    _fig.tight_layout()
    plt.show()


plot_latent_tsne(
    prot_z, protein_vqvae.codebook, config.protein.codebook_size, "Protein"
)
plot_latent_tsne(
    lig_z, ligand_vqvae.codebook, config.ligand.codebook_size, "Ligand"
)

```

Protein — t-SNE of latent space



Ligand — t-SNE of latent space



7. Summary table

```
In [ ]: import pandas as pd

prot_idx_np = prot_indices.cpu().numpy()
lig_idx_np = lig_indices.cpu().numpy()

prot_probs = np.bincount(
    prot_idx_np, minlength=config.protein.codebook_size
).astype(float)
prot_probs = prot_probs / prot_probs.sum()
prot_probs_nz = prot_probs[prot_probs > 0]
prot_ppl = np.exp(-np.sum(prot_probs_nz * np.log(prot_probs_nz)))

lig_probs = np.bincount(lig_idx_np, minlength=config.ligand.codebook_size).a
    float
)
lig_probs = lig_probs / lig_probs.sum()
lig_probs_nz = lig_probs[lig_probs > 0]
lig_ppl = np.exp(-np.sum(lig_probs_nz * np.log(lig_probs_nz)))

summary = pd.DataFrame(
    {
        "Metric": [
            "Codebook size",
            "Latent dim",
            "Test MSE (normalized)",
            "Test MSE (original scale)",
            "Test RMSE (original scale)",
            "Active codes",
            "Utilization",
            "Perplexity",
            "Perplexity (normalized)",
        ],
        "Protein": [
            config.protein.codebook_size,
            config.protein.latent_dim,
            f"{prot_metrics['mse_norm']:.6f}",
            f"{prot_metrics['mse_orig']:.6f}",
            f"{np.sqrt(prot_metrics['mse_orig']):.6f}",
            len(np.unique(prot_idx_np)),
            f"{len(np.unique(prot_idx_np)) / config.protein.codebook_size:.4f}",
            f"{prot_ppl:.1f}",
            f"{prot_ppl / config.protein.codebook_size:.4f}",
        ],
        "Ligand": [
            config.ligand.codebook_size,
            config.ligand.latent_dim,
            f"{lig_metrics['mse_norm']:.6f}",
            f"{lig_metrics['mse_orig']:.6f}",
            f"{np.sqrt(lig_metrics['mse_orig']):.6f}",
            len(np.unique(lig_idx_np)),
            f"{len(np.unique(lig_idx_np)) / config.ligand.codebook_size:.4f}",
            f"{lig_ppl:.1f}",
            f"{lig_ppl / config.ligand.codebook_size:.4f}"
        ]
    })
```

```

        f"{lig_ppl / config.ligand.codebook_size:.4f}",
    ],
}
)
summary

```

	Metric	Protein	Ligand
0	Codebook size	2048	1024
1	Latent dim	16	8
2	Test MSE (normalized)	0.130986	0.017311
3	Test MSE (original scale)	0.084838	0.010141
4	Test RMSE (original scale)	0.291270	0.100704
5	Active codes	2028	1024
6	Utilization	0.9902	1.0000
7	Perplexity	1594.4	763.1
8	Perplexity (normalized)	0.7785	0.7452